

Data Mining: Learning from Large Data Sets - Spring Semester 2014

Project 1

nivo@student.ethz.ch
ganzm@student.ethz.ch
rohrp@student.ethz.ch

March 12, 2014

Approximate near-duplicate search using Locality Sensitive Hashing

This document explains the implementation of project 1. In this first project we were supposed to find near-duplicate videos using Locality Sensitive Hashing. The videos were given as a number of shingles that could be compared. The implementation was done in Python to be used in combination with the hadoop infrastructure. This means we implemented a mapper and a reducer that read from stdin and write to stdout. In the following two sections we explain the functionality of the mapper and the reducer each.

mapper.py

The mapper creates a signature vector of a video. This signature vector is afterwards subdivided into b bands having r elements that get hashed. Each band and its hash are then emitted.

The signature is calculated using min-hashing:

1. As Figure 1 shows, choosing $b = 32$ and $r = 8$ yields no true positives.
2. We create $k = b \cdot r$ hashfunctions of the form $a_p x + b_p \mod c_p$ where $a_p = rand(0, 999)$, $b_p = rand(0, 99999)$, $c_p = 10000$ and $rand(x, y)$ is random variable drawn from a uniform distribution between x and y . These hashfunctions are used to calculate the permutations in min-hashing.
3. We create r hashfunctions of the form $a_b x + b_b \mod c_b$ where $a_b = rand(0, 999)$, $b_b = rand(0, 999)$, $c_b = 1000$ and $rand(x, y)$ is random variable drawn from a uniform distribution between x and y . These hashfunctions are used to calculate the band hashes before emitting.
4. The signature is then created according to 1
5. Calculate the band hash and emit the band hash and the band as key, and the video with its corresponding shingles as value according to 2




Figure 1: $1 - (1 - x^r)^b$

Algorithm 1 Create signature

```

signature = ∞
for all shingle in shingles do
  for i in range(k) do
    signature[i] = min( $a_{p_i} \cdot shingle + b_{p_i} \bmod c_p$ , signature[i])
  end for
end for

```

Algorithm 2 Emit keys and values

```

for band in range(b) do
  vector = signature[band · r : band · r + r]
  bandhash =  $\sum_{i=1}^{len(vector)} a_{b_i} \cdot vector[i] + b_{b_i} \bmod c$ 
  emit key = [bandhash, band] value = [videoid, shingles]
end for

```

reducer.py

The main task of the reducer is to get rid of the false positives by comparing the reported similar videos using the jaccard distance:

1. gather all videos with the same keys in a collection *duplicates*
2. emit similar videos like shown in 3

Algorithm 3 Emit similar videos

```
for i=0 to len(duplicates) do
  for j=i+1 to len(duplicates) do
    if duplicates[i].video_id < duplicates[j].video_id then
      shingles_left = duplicates[i].shingles
      shingles_right = duplicates[j].shingles
       $distance = \frac{|shingles\_left \cap shingles\_right|}{|shingles\_left \cup shingles\_right|}$ 
      if distance > 0.85 then
        emit duplicates[i].video_id duplicates[j].video_id
      end if
    end if
  end for
end for
```
